

Modern Mixed-Precision Methods in Portable C++ for Accelerated Hardware Platforms

Neil Lindquist

SIAM PP22

February 26, 2021



Modern Numerical Methods

Precision

- Double
- Single
- Half
- Bfloat16
- Integer?
- Posits?

Hardware

- CPU
- GPUS
 - NVIDIA
 - AMD
 - Intel
- ML Accelerators
- FPGAs

Mixed Precision Experiments

- Many combinations of precisions to test
 - My GMRES work: 14 combinations in paper
 - Only single, double
 - GMRES-IR5¹: 112 “meaningful” combinations
 - Bfloat16, half, single, double, quad

Possible Types of Configuration

1. Read time
2. Compile time
3. Run time

Read time configuration

- Limited to string replacements
 - E.g., macros
- Error prone

Possible Types of Configuration

- ~~1. Read time~~
2. Compile time
3. Run time

Compile time configuration

- Templating - Parameterize code
 - Both types and constants

Templating Type

```
template<class T>
T dot(int n, T* x, T* y) {
    T sum = 0.0;
    for (int i = 0; i < n; i++) {
        sum += x[i]*y[i];
    }
    return sum
}
```


Compile time configuration

- Templating - Parameterize code
 - Both types and constants
- Can specialize for specific configurations
 - Different vendor libraries
 - Host vs device memory

Templating Devices

```
template<class T, class Target>  
T dot(Vect<T, Target> x,  
      Vect<T, Target> y);
```

```
template <>  
double dot<double, MKL>(Vect<double, MKL> x,  
                        Vect<double, MKL> y) {  
    return cblas_ddot(x.n(), x.data(), 1,  
                      y.data(), 1);  
}
```

Templating Devices

```
template <>
double dot<double, Cuda>(Vect<double, Cuda> x,
                        Vect<double, Cuda> y) {
    double result;
    cublasDdot(cublas_handle, x.n(),
               x.data(), 1,
               y.data(), 1,
               &result);
    return result;
}
```

Kokkos

- Performance Portability Library
- View – generic multi-dimensional array
 - Template: type, device, memory layout
- Device-portable kernels
- Similar ideas in Raja, ect.

Templating with Kokkos

```
template<class T, class S, class MemSpace,
        class R=std::common_type<T, S>>
R dot(Kokkos::view<T*,MemSpace> x, Kokkos::view<S*,MemSpace> y){
    R out; int n = x.n();
    Kokkos::parallel_reduce(
        Kokkos::RangePolicy<typename MemSpace::execution_space>(0,n),
        KOKKOS_LAMBDA(int i, R& partial_sum) {
            partial_sum += x(i)*y(i);
        }, Kokkos::Sum(out));
    return out;
}
```

Possible Types of Configuration

- ~~1. Read time~~
2. Compile time
3. Run time

Run time configuration

- More flexible APIs
- Manually wrap templated code

Wrapping Templated Routines

```
double dot(std::string target,  
           Vect<double> x, Vect<double> y) {  
    if (target == "MKL") {  
        return dot<MKL>(x, y);  
    }  
    if (target == "CUDA") {  
        return dot<CUDA>(x, y);  
    }  
    ...  
}
```

Wrapping Templated Routines

```
Scalar dot(std::string target,  
           std::string x_type, Vect x,  
           std::string y_type, Vect y);
```

- 2 devices & 2 precisions → 8 branches
- 2 devices & 3 precisions → 18 branches
- 3 devices & 4 precisions → 48 branches

Run time configuration

- More flexible APIs
- Can manually wrap templated code
- Runtime types uncommon
- DPC++: runtime devices
- OCCA: runtime types, devices

Conclusions

- Combinatorial explosions of implementations
 - Precisions and accelerators
- Can alleviate with
 - Templating
 - Portability libraries